

Q1

1.
 - a. 3. One for each host.
 - b. Since UDP does not make distinctions between different ports, the socket on S1 will only have 1 socket for host B. Therefore, the packet will be sent to that socket.
 - c. Just one. Each host makes a single socket to connect to a single server, but since the applications are communicating to the same server, each host needs only 1.
 - d. The packet will get sent to B2's only socket that interfaces with S1. Since UDP has a single socket for all S1 connections, the packet will reach that socket and the correct application will read the packet.
2.
 - a. 6. Each port and IP needs a new socket.
 - b. The packet will be sent to the socket with the specific port and B2's IP on S1.
 - c. A will need 1, B will need 3, and C will need 2. Each host needs a separate socket for each application because each application uses a different port.
 - d. The packet will be sent to a specific socket that corresponds to the connection between B2 and S1 with a specific port.

Q2

1. TCP - very reliable. Banking apps can't have dropped packets.
2. Hmm. UDP would be ok if there wasn't a focus on reliability. Since there are so many devices it would be nice if you didn't have to make a new socket for each port.
3. UDP - needs a large amount of data delivered fast. If some data is dropped the user will likely be able to infer what the missing information is.
4. UDP - needs a large amount of throughput and users do not expect perfect audio quality.

Q3

1. A needs to wait for an ACK from B every time.
100Mbps = around 100000000bps
So three packets =
 $3 * 2 * (1000\text{bits}/100000000 \text{ bps} + 3\text{ms}) = \mathbf{18.06\text{ms}}$
Throughput = $3000\text{bits}/18.06\text{ms} = \mathbf{166,113 \text{ bps}}$
2. A sends 3 packets to B, then waits for 3 acks.
 $(1000\text{bits}/100000000 \text{ bps}) * 2 + 3\text{ms} * 2 = \mathbf{6.02\text{ms}}$
Throughput = $3000\text{bits}/6.02\text{ms} = \mathbf{498,339 \text{ bps}}$

Q4

1.
 - a. 0
 - b. There will be a timeout where A doesn't see Ack1.
 - c. It will send packet 1, 2, 3.
2.
 - a. 2
 - b. There will be a timeout when A doesn't see Ack1.
 - c. It will only send packet 2.

Q5

1. They're all ack146 because A never resends packet 146!
2.
 - a. After the timer for packet 146 expires
 - b. Assuming window size of 6: 146, 162, 178, 196, 212, whatever packet is next
 - c. Ack162
3.
 - a. After the 3rd duplicate ack146
 - b. Just packet 146
 - c. Ack220, since the rest of the packets delivered correctly

Q6

1. 8. After cwnd reaches cwnd=8, the graph switches from slow start (the exponential growth) to additive increase (the linear sections).
2. There was a packet loss (3 duplicate acks)! After that happens, Reno specifies that cwnd should return to ssthresh, which in this case is 8.
3. There was another packet loss due to 3 duplicate acks. Ssthresh is set to 5 because the packet was lost at cwnd=11, and $11/2=5$. Fast recovery specifies that cwnd should return to ssthresh=5.
4. A packet was lost due to timeout, causing the drop to 1.
5. Ssthresh=6 because the last time a packet was lost was when cwnd=12, so $12/2=6$. Therefore, additive increase is used. $8+1=9$.